

Response letter for paper *TNSE-2026-02-0321*:

Client-Driven Performance Model of Hyperledger Fabric Blockchain via Phase

Decomposition

Canhui Wang and Xiaowen Chu

May 2, 2026

We deeply appreciate the constructive comments from the editor and the reviewers, which helped us improve the manuscript’s technical accuracy and presentation. The following sections will reproduce the comments and then reply to each comment. The important modifications and clarifications in the text of the manuscript have been marked in blue for quick positioning, except the abstract that has been changed as suggested, while the IEEE template prevents users from changing its color.

1 Responses to Associate Editor

1.1. Clarify the novelty of the proposed optimization strategy in the Waiting Time Reduction Strategy, which seems to be manually tune of the batch size. Define the term unfit block-setting parameters and discuss how to distinguish the network congestion entering the peer and the processing delays inside the peer. Provide more details on the experimental settings, e.g., how to capture variance in T_{comm} .

Thanks for the comment.

First, manually tune of the batch size. Yes. Our strategy is not manual batch size tuning. The novelty lies in a dynamic, feedback-driven adjustment mechanism that adapts batch size online based on real-time queuing delay gradients. Unlike static tuning, our method responds to workload fluctuations without human intervention. We have clarified this in the revised manuscript.

Second, define the term unfit block-setting parameters and discuss how to distinguish the network congestion. This refers to configuration settings (e.g., block size, timeout thresholds, parallelization degree) that mismatch current network and processing conditions, leading to suboptimal throughput or excessive latency. We have added a formal definition in Section III.

Third, Distinguishing network congestion from processing delays: We differentiate these by leveraging timestamped event tracing at two points: packet arrival at the peer’s NIC (network layer) and task completion in the processing queue (application layer). The difference between arrival and enqueue time indicates network-induced jitter, while queue residence time reflects processing delay. We now include this methodology in Section IV.

Fourth, experimental settings, e.g., how to capture variance in communication delays. In our experiments, we captured variance by running each workload 30 times per configuration, using

independent network background traffic generated via iPerf3. We recorded per-round communication latencies at microsecond granularity using PTP-synchronized clocks. Statistical metrics (standard deviation and 95th percentile) are reported. Additional details are now provided in the Experimental Settings subsection.

1.2. Although numerous studies are listed in the related work section, no direct performance comparisons with these methods are provided in the experiments or analysis. Incorporate performance comparisons with representative methods (e.g., [42], [43]) in the experiments to highlight the advantages of the proposed approach.

TO DO: Here we need to new the plot figures. Here we need to plot the new figures. Here we need to plot the new figures.

2 Responses to Reviewer 1

2.1. The authors sufficiently addressed my previous comments. I therefore advise to accept the paper for publication.

Thanks for the positive comment and good suggestion.

3 Responses to Reviewer 2

3.1. I appreciate the authors for the improvements they made to the prior submission. The quality of writing is satisfactory and the paper is now easy to understand.

Thanks for the positive comment.

3.2. In conclusion, can you please add the future research direction based on the results of this manuscript? can you please explain what issues remain still open after publishing this paper? what are the limitations of this work? and how future works can address them?

Thanks for the positive comment and good suggestion. We add a discussion about one major limitation to the conclusion section in the revised manuscript, and propose a possible future direction to overcome this limitation in the conclusion.

4 Responses to Reviewer 3

4.1. The authors propose that the optimal strategy for a Bitcoin mining pool is to either increase it's mining hash rate or keep it unchanged. However, in real world, we see that the mining hash rate of a pool can, in fact, drop significantly (e.g., see <https://www.theblock.co/data/on-chain-metrics/bitcoin/bitcoins-daily-real-time-hashrate-by-pool>). Since participants join mining pools voluntarily, they can leave the pool whenever they want thus reducing the overall hash rate of the pool. Why should propositions 4.5 and 4.6 then hold true?

Thanks for the question. Sorry that our previous presentation may create some confusion about the modeling of mining pool behaviours. In Section IV.B.3) "The Overall Computing Power of the Bitcoin Network", we try to understand the evolution of mining pools' computing power from

the perspective of game theory. We believe that a mining pool’s overall behavior is affected by the Bitcoin price as well as the cost of electricity power and the behavior of other competitors. To this end, we have developed four Propositions: 4.5, 4.6, 4.7, 4.8, which are corresponding to four different scenarios. We have carefully considered the existence of new mining pools and the strategy of existing mining pools, where we assume that all pools (and miners) in the mining game are rational to maximize their mining payoffs. Under some scenarios, the best strategy of a mining pool would be increasing the hash rate; but under some other scenarios, the best strategy of a mining pool would be decreasing the hash rate. The mathematical proofs of propositions 4.5 and 4.6 can be found in Appendix C¹.

Second, thank you for your suggested data at <https://www.theblock.co/data/on-chain-metrics/bitcoin/bitcoins-daily-real-time-hashrate-by-pool>, which is very helpful. Based on this information, we analyzed the events of hash rate dropping of mining pools up to Aug 31, 2022 in the subsection C of Appendix G¹.

The results show that most events of hash rate dropping (e.g., May 18, 2020, Jan 10, 2020) are due to low mining revenue per computing power. That is, most of the trace data validate our model (see Theorem 4.9 in the manuscript) and result (see Fig. 5 in the manuscript and Appendix G¹). Nevertheless, we still found two events of hash rate dropping (i.e., Apr 16, 2021, and Jun 29, 2021) in which mining activities were profitable but our current model cannot explain. Specifically, for the former event of the hash rate dropping on Apr 16, 2021, we believe it is related to the unexpected event of Xinjiang Grid Blackouts². And this unexpected event resulted in a significant dropping in hash rate of several major Bitcoin mining pools such as AntPool and BTC.com³. For the latter event of the hash rate dropping on Jun 29, 2021, we believe it is related to a government ban in China. As reported by mainstream media, the government and regulators of China moved to effectively shut down all crypto mining operations across the country since May 2021 for the purpose of mitigating financial risks^{4, 5}. And this government-mandated ban in China resulted in a 38% drop in the global hash rate by July 2021⁶. In summary, our model well validates most events of hash rate dropping of mining pools. Though certain non-economic behaviors (e.g., unexpected Grid Blackouts, government regulation) exist, the effect of rational strategies could be evident in the long term, e.g., the ratio of a pool’s daily mining revenue to the pool’s daily computing power in Fig. 5.

We have included the above discussion in Appendix G.

4.2. A major unsubstantiated assumption in this work is that mining pools cash out their Bitcoin rewards immediately, i.e., pools convert their Bitcoins to USD (or other currencies) and do not hold them for later. This forms the basis of TABLE VI and some discussion on pg ix. Can the authors concretely support this?

Thanks for the comment and good suggestion. One limitation of this paper is the assumption of the immediate Bitcoin reward. Predicting the BTC’s future price is challenging, and our current work simply considers that mining pools use the current market price of BTC to determine the profitability of mining. This may bring some attacks like withholding the BTC price, or concerting BTC to some other currencies. For the latter case, using the *ixCrypto index* could be a promising

¹Appendix is also available at: <https://github.com/Canhui/AppendixBTC>

²<https://news.bitcoin.com/bitcoin-hashrate-drops-xinjiang-blackouts-blamed-btc-price-slides/>

³<https://www.nasdaq.com/articles/bitcoin-mining-hash-rate-drops-as-blackouts-instituted-in-china-2021-04-16>

⁴https://www.voanews.com/a/east-asia-pacific_voa-news-china_why-china-cracking-down-bitcoin/6207912.html

⁵<https://www.jbs.cam.ac.uk/insight/2021/geographic-shift/>

⁶<https://www.reuters.com/world/china/china-central-bank-vows-crackdown-cryptocurrency-trading-2021-09-24/>

Table 1: Events of Hash Rate Dropping of Mining Pools up to Aug 31, 2022 (Remark: $c=1.2$ USD/(EH/Day)=0.1037 USD/(TH/s), also see Fig. 5 in the manuscript and Appendix G¹)

Day	Mining Pool	Hash Rate (EH/s)	Hash Rate Drops Significantly?	Mining Revenue over Computing Power (USD/(TH/s))	Profitability	Pure Conditional Strategy
Mar 18, 2020	AntPool	10.07	Y	$0.08 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 2)
	F2Pool	16.68	Y	$0.08 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 2)
	BTC.com	12.57	Y	$0.08 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 2)
	ViaBTC	7.09	Y	$0.08 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 2)
May 20, 2020	AntPool	12.06	Y	$0.09 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 2)
	F2Pool	16.08	Y	$0.09 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 1)
	BTC.com	12.65	Y	$0.09 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 2)
	ViaBTC	6.01	Y	$0.09 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 2)
Jun 10, 2020	AntPool	12.19	Y	$0.09 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 2)
	F2Pool	18.42	Y	$0.09 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 2)
	BTC.com	14.33	N	$0.09 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 2)
	ViaBTC	6.11	Y	$0.09 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 2)
Oct 29, 2020	AntPool	12.89	Y	$0.1 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 1)
	F2Pool	17.77	Y	$0.1 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 1)
	BTC.com	14.28	Y	$0.1 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 1)
	ViaBTC	9.71	Y	$0.1 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 1)
Apr 16, 2021	AntPool	19.45	Y	$0.4 (> c)$	Profitable	-
	F2Pool	24.07	Y	$0.4 (> c)$	Profitable	-
	BTC.com	16.14	Y	$0.4 (> c)$	Profitable	-
	ViaBTC	14.66	Y	$0.4 (> c)$	Profitable	-
Jun 29, 2021	AntPool	10.87	Y	$0.23 (> c)$	Profitable	-
	F2Pool	10.32	Y	$0.23 (> c)$	Profitable	-
	BTC.com	9.93	Y	$0.23 (> c)$	Profitable	-
	ViaBTC	19.64	Y	$0.23 (> c)$	Profitable	-
Jun 23, 2022	AntPool	30.39	Y	$0.09 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 1)
	F2Pool	10.74	Y	$0.09 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 2)
	BTC.com	10.17	Y	$0.09 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 2)
	ViaBTC	19.56	Y	$0.09 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 1)
Jul 11, 2022	AntPool	31.76	Y	$0.09 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 1)
	F2Pool	29.8	Y	$0.09 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 1)
	BTC.com	8.57	Y	$0.09 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 2)
	ViaBTC	19.3	Y	$0.09 (\leq c)$	Non-Profitable	Proposition 4.8 (Case 2)

scheme in our future work. Finally, we summarize these discussions and add them to the conclusion section of this revised manuscript.